

Cognitia — Session & Project Audit

Generated 2026-06-16 from the durable cross-session record (git/GitHub + the local Claude Code profile). "Sessions" = the Claude Code work/co-work sessions run on this profile; each `claude/*` branch $\approx$ one session.

1. Executive summary

- **Repository:** `cognitiacloud/cognitia.cloud` · default branch `main` @ `fback05` .
- **Sessions reconstructed:** $\sim$ **90** `claude/*` **branches** $\rightarrow$ **87 pull requests**.
- **PR status:** **69 merged** · **9 open (in flight)** · **9 closed-unmerged (abandoned/superseded)**.
- **Active span:** 2026-05-28 $\rightarrow$ 2026-06-16 ($\sim$ 3 weeks of near-daily sessions).
- **Two work threads:**
 1. **Cognitia — Governed GTM Action System** (the product; the bulk of the work).
 2. **Cognitia media / Hermes content pipeline** (earlier; mostly superseded; the surviving artifact is `hermes/skills/vision-skill/`).
- **External trackers** (Linear / Notion / Airtable): effectively **empty** — no Cognitia projects/tasks are tracked outside GitHub (see §5).

2. What the system is

Cognitia — a Governed GTM Action System. A TypeScript-first, **CRM-first, approval-gated** GTM automation product. The live agent **Mira** proposes HubSpot CRM writes — two governed write types today: **follow-up tasks** and **grounded account-context notes**. Every side effect runs a governed lifecycle:

preview $\rightarrow$ human approval (mandatory structured reason) $\rightarrow$ idempotent, provenance-stamped execution $\rightarrow$ reversible undo

...and every capability claim is backed by a test that fails CI on regression.

- **Monorepo (pnpm + TypeScript):** `apps/web` (Next.js operator console), `apps/api` (Fastify-style service), `apps/worker` (jobs); `packages/` `core` · `db` · `agents` · `integrations` · `evals` · `workflows` .
- **Stack:** Supabase Postgres + pgvector · Zod · Vitest · n8n workflows.
- **Agents:** Mira = **Live (v1)**; Echo / Atlas / Beacon = Planned/Later.
- **Scope fence (V1):** CRM write-back only (HubSpot tasks/notes). Email, voice, ads execution are **fenced off** and enforced in tests + deploy smoke.

Second thread — Cognitia media / Hermes pipeline: video "Episode 002" (Remotion scaffold), HeyGen, founder-avatar vision QC, and a Windows $\leftrightarrow$ Hermes filesystem mesh bridge. Surviving artifact: `hermes/skills/vision-skill/` — read-only local OCR + multi-provider vision QC (OpenAI/Anthropic/Gemini/ OpenRouter/Ollama, OCR-only fallback), 13 passing tests, CLI + MCP stdio server.

3. Session timeline (oldest $\rightarrow$ newest)

Date	Session group	PRs	Outcome
------	---------------	-----	---------

May 28–30	Media / Hermes thread — hermes vision skill; Episode 002 rebuild (Remotion); Windows↔Hermes bridge	commit + #1, #2	skill kept; #1 #2 closed-unmerged
Jun 6	GTM Phase 0 foundation + Mira outbound MVP	#3	open
Jun 10	Governance wave — decision reasons, provenance, batch UX, trust metrics, eval gate, typed write/preview, preflight, undo, trust packet, regression flywheel, kill switch, readiness gate, lifecycle acceptance, rationale, README coherence, scorecards, grounded notes, run/plan review, hardening, SOC	#4–#30 (excl. #18)	all merged
Jun 11	Cognitia v1.1 / COG pack — schema foundation, proof registry, Agent Trust Credential, SkillProof + AI Front Desk/Lead Rescue, reputation v0, credits/wallet, command dashboard, evidence sync, run timeline, capability ledger, ally smokes	#32–#43 (excl. #31)	all merged
Jun 12–13	Agent Economy + security — SEC-1, Agent-Economy 001–005 (lab, disputes, agent actions, marketplace, settlement), agent passports, token architecture (doc), economy smoke, mainline status	#47–#57 (+ #50)	merged; #44 #45 #46 #54 open
Jun 14	Public trust / visibility — crypto-visibility, Trust/Proof Explorer V-4/4b/4c, V-5 feed hardening, 12h research, researcher pack, diligence discoverability	#58–#66	all merged
Jun 15	Visibility + Lanes + Audit Booklet — API reference, threat model, Agent Fabric Lab, system booklet, Lanes C/D/E (HubSpot live, AI drafting, meeting intel), overnight orchestrator, PILOT-001 mainline harness	#67–#83 (subset)	mostly merged; #71–74, #82 closed
Jun 15–16	RLS hardening — managed Postgres RLS harness, PILOT status board, V6A docs reconcile	#84, #85, #87	merged
Jun 16	This session — claude/plot-sessions-audit (this audit)	—	in progress

4. Status board

✅ Done — shipped to main (~69 merged PRs)

The full governed CRM loop is **live and CI-proven**: per-action approval with a mandatory reason; "why this action" rationale + data-freshness warning; byte-identical write preview; zero-write preflight simulation; connection-readiness gate; idempotent provenance-stamped execute; accountable undo; tenant kill switch; code-derived governance matrix + audit trail; trust metrics + exportable trust packet; run/plan review; rejection→regression flywheel; falsifiable golden eval gate; full-lifecycle acceptance test; post-deploy smoke; **Postgres RLS tenant isolation**. Plus: public-safe /trust Proof Explorer + live proof feed + diligence docs (researcher pack, threat model, API reference);

Agent Economy Lab (simulation-only: work orders, escrow, disputes, marketplace, settlement design); token architecture spec (doc-only); the **AUDIT-BOOKLET** master system booklet.

 In flight — 9 open PRs

PR	Title	Since
#3	Phase 0 foundation + Mira outbound MVP	Jun 6
#44	COG-011: Lead detail — story endpoint + console	Jun 12
#45	COG-011 + COG-012: Lead detail + tenant provisioning	Jun 12
#46	COG-014: Demandara onboarding (2nd vertical GTM loop)	Jun 12
#54	Agent-Economy-004: marketplace listings + matching	Jun 12
#61	Fix Hermes bridge MCP stdio restart loop	Jun 14
#78	Lane B: operator Approval Queue + Run visibility	Jun 15
#79	COG-011 (UI): Lead Detail Console page	Jun 15
#86	Meeting-notes writeback via governed <code>crm.note</code>	Jun 15

 Abandoned / superseded — 9 closed-unmerged PRs

#1 Episode 002 rebuild · #2 Windows Hermes bridge · #18 agent-economy 2-week spec · #31 COG-001 discovery · #71 CODE 28→50 integration manifest · #72 Lane A operator-UI shell · #73 LEGEND-001 stitch/merge · #74 Lane C v1 (→ superseded by merged #77) · #82 PILOT-001 Tenant Zero (→ superseded by merged #83).

Known remaining work (not unbuilt — human/data-blocked)

- **Live operator setup** needs human-provided HubSpot credentials + portal setup.
- **Earned-autonomy / risk-tiered review** is gated on accumulated decision-label volume.

5. Tracked elsewhere? (external scan)

You asked to include external trackers. Result — **no Cognitia work lives outside GitHub:**

- **Notion:** search for "Cognitia" → 0 results.
- **Airtable:** `list_bases` → no bases.
- **Linear:** a team "**Cognitiacloud**" exists but has **no projects** and only the 4 default onboarding issues (COG-1...COG-4 : "Get familiar with Linear", etc.). Note: Linear issue keys COG-* are unrelated to the repo's COG-00x work items.

6. Methodology & caveats

- **Sources:** local Claude Code profile (`~/ .claude/projects/`), GitHub branches / PRs / commits for `cognitiacloud/cognitia.cloud` , and the Linear/Notion/Airtable MCP integrations.
- **Ephemeral transcripts:** web/remote Claude Code containers are reclaimed after use, so only the *current* session's transcript persists locally. Past sessions are reconstructed from their

durable git footprint (branch + PR), not transcripts — so per-session prompt detail is not recoverable, only the work product.

- **"Plot / co-work / court sessions"** are read as the Claude Code work sessions.
- **PR counts** (87 / 69 / 9 / 9) reflect state at generation time; re-run `list_pull_requests` to refresh.