

Cognitia / Demandara -- Master Execution Report (Latest)

Date: 2026-06-23 Canonical line: overnight/gtm-implementation @ da48e8f1beeb2709591e7951d49fa3a893cb4d47 (Merge PR #179, "pure automation release-gate engine"). Compiled by: master-report session on branch claude/cognitia-demandara-master-report-t89yho (rebuilt from the canonical line; docs-only).

Mandate: one honest, citable reconciliation of the whole line. No fake proof. No inflated scores. Actual-live is capped by design. Every number cites a file, test, route, command, or PR.

Evidence legend

Mark	Meaning
[x]	verified Re-run or read directly on da48e8f in this session (command output / file / test)
claim (PR)	Asserted by an open PR or branch doc; NOT yet on the canonical base
simulated	Real code, but mock / dry-run / sandbox by construction -- never live
[?]	not found Named input searched for and not present in the repo (not fabricated)

0. This session's verification (the only first-party proof here)

Run on the canonical tip da48e8f after pnpm install --frozen-lockfile (the CI install, .github/workflows/ci.yml):

Command	Result	Sta
pnpm run check (format:ch exit 0 prettier clean tsc clean	805 tests passed / 106 files (39.2s)	[x]
pnpm --filter @cognitia/w exit 0 ~19 routes render (/cognitia, /credits, /proofs, /skills, /trust, /		[x]

The 805/106 figure is today's canonical base, and supersedes the higher counts quoted by individual audit PRs (786 in #158, 771 in the integration plan, 829 in #164/#181) -- those scored different, not-yet-merged consolidated branches at earlier moments. Where this report cites those numbers it labels them claim (PR).

1. What is built ([x] on the canonical base)

Cognitia is a pnpm/TypeScript monorepo -- "AI GTM workforce platform" (package.json). Three apps over six packages, all building and tested on da48e8f.

Apps

- apps/api -- Fastify v5 server (apps/api/src/server.ts, buildServer()). ~77 session-authed routes defined server.ts:131-481. Tenant is derived from the session principal, never from request headers (apps/api/src/auth.test.ts). [x]
- apps/web -- Next.js 15 / React 19 app-router operator console (apps/web/src/app*/page.tsx). next build renders ~19 routes. [x]
- apps/worker -- background-job scaffold (apps/worker/src/index.ts); runRegisteredJobs() shell, no active jobs yet. scaffold.

Packages (packages/*, names from each package.json)

- @cognitia/core -- schemas, events, policies, logging, doctrine guards. [x]
- @cognitia/db -- Kysely + in-memory twin repository; SQL migrations 0001-0019; RLS. [x]
- @cognitia/agents -- Mira (scoring/drafting), GTM-OS assembly, dry-run channels, closer, economy, fabric, TrustOps. [x]
- @cognitia/integrations -- HubSpot client (real OAuth in prod / fake in dev), webhook signature verify, credential cipherng. [x]
- @cognitia/evals -- golden datasets + regression harness. [x]
- @cognitia/workflows -- placeholder (src/index.ts empty). stub.

The governed product loop (closed and tested): preflight (zero-write) -> propose (evidence+policy) -> preview (byte-equal to write, CI invariant) -> approve (mandatory reason) -> execute (idempotent, provenance-stamped) -> undo (reversible) -> reject feeds the regression gate -> exportable trust packet. Each stage cites a test: previewAction.test.ts, preflight.test.ts, rollback.test.ts, regressionCandidate.test.ts, trustPacket.ts. [x]

Surfaces that build & render (apps/web/src/app): /, /approvals (operator queue), /cognitia (COG-007 command dashboard), /proofs, /agents + /agents/[id], /skills, /credits, /cognitia/crypto-readiness, /moveros/front-desk, /discovery, /trust, /trust/live, /gtm-os-integrated-demo, /portal/*. [x]

Key built capabilities, each tested: approval-gated action ledger; enforced kill-switch (apps/api/src/killSwitch.test.ts -- non-active connection -> 409 + audited denial); reversible rollback; real HubSpot read-sync + signature-verified ingest webhook (fail-closed, apps/api/src/webhookHubspot.test.ts); Proof Registry (append-only, redaction-gated); ATC lifecycle; SkillProof; verified-fact-only Reputation; internal double-entry Credits ledger + inert wallet placeholders; Agent Economy work-order/marketplace loop (simulation); Agent Fabric Lab (simulation); per-segment scorecards. [x] / as marked in 6 and 13.

2. What is canonical (merged on overnight/gtm-implementation @ da48e8f)

Canonical = what is actually on the audited base, not main. Divergence ([x] git rev-list): overnight/gtm-implementation is 10 commits ahead of main (d3d198e7), and main is 0 commits ahead -- i.e. main is strictly behind and fully contained in the canonical line (merge-base = main's head). main has not caught up; do not treat it as canonical. [x]

Canonical content = everything in 1 above, plus the merged closer-automation primitives: deterministic channel policy + dry-run channels (packages/agents/src/channels/{channelPolicy,dryRunChannels}.ts), the kill-switch surface, and the automation release-gate engine (automationReleaseGate.ts, merged as PR #179 = the da48e8f tip). [x]

The 2026-06-10 "Top-10" governance roadmap is substantially shipped on base (6, Top-10 table): 7 shipped with tests, 1 partial, 2 open by design -- per docs/strategy/roadmap-audit-2026-06-23.md (PR #188), reconciled against code/test paths on da48e8f. [x] /

Not yet canonical (important): the /gtm-command-center route does not exist on the base -- it lives only in branch PRs (#160/#161/#167/#168/#169). The web build on da48e8f does not emit it. [x]

3. What is branch-only (-- real work, not on the base)

Per the GitHub API on 2026-06-23: ~40 open PRs against this line (99 open repo-wide), ~98 drafts, 0 failing CI, but deeply stacked. The canonical line itself is open PR #158, stacked ~3 levels from main. This is the program's binding constraint (8, 12). Source: docs/strategy/roadmap-audit-2026-06-23.md 3.

Group B -- the active line (22 PRs targeting overnight/gtm-implementation):

Theme	PRs	Note
GTM Command Center (>=5 compe	#160, #161, #167, #168, #169, #186/#189	One must be chosen canonical; res
Integration packet (B1-B6 uni	#159	Base of the Command Center work
Alta 80/90 readiness audits	#164, #178, #181	Consolidate to one
Proof / TrustOps evidence (du	#165, #166	Near-identical; pick one
Closer automation-readiness s	#171->#175->#176->#174->#173->#177->#18	Cleanest sub-stack; #179 release-
Enterprise-readiness controls	#162, #185	Mock-safe
Investor panel	#190	Read-only, dry-run
Roadmap audit (this date)	#188	The 99-PR rollup this section dra

Branch-only audit docs (cited in 6/13 as): alta-90-final-readiness-evidence.md (PR #164), live-automation-80-readiness-evidence.md (PR #181). They are not on the base; the base carries the earlier docs/cognitia/audits/alta-80-readiness-evidence.md. [x] (base) / (branch docs)

4. What is wrong-repo / noise / superseded

- The stranded vision-skill branch line. This report's own branch was previously parked at 0dfb0ad ("Add hermes vision skill") -- a single commit detached from the product line. Branch claude/ep002-mission-run-pPoba

carries that line and acts as an integration branch with ~24 PRs feeding it but no PR of its own (roadmap-audit 3.4); none of those are mainline-reachable. Treat the vision-skill-only state as out of the canonical GTM line (it is a Hermes QC skill, not GTM code). [x]

- Duplicate / competing PRs (convergence tax, not new value): the >=5 Command Center attempts and the #165/#166 TrustOps pair (3). These are "noise" only in the sense that all-but-one should be closed.
- main as a reference point -- strictly behind the line (2); citing main for "what's built" would understate reality. Use the canonical line.
- No fabricated inputs. Several named inputs do not exist in the repo and are recorded as [?] in 5 -- they are not invented here.

5. What is missing (gaps + named inputs not found)

Named inputs -- search results ([?] = not fabricated):

Requested input	Status	Closest real artifact
"Cognitia Republi	[?] not found	-- (no "Republic" reference in repo)
"Trust From Zero"	[?] not found	b docs/competitive/operating-plan.md 8 (30/60/90 execution); docs/strate
"War Council" str	[?] not found	b docs/cognitia/research/12H_CRYPT0_VISIBILITY_AGENT_FABRIC/FOUNDER_COUN
Alta / SalesClose	[x] found	docs/strategy/beat-alta-10x.md, docs/AltaSpec_v2.yaml, docs/sales-clos
Agent Economy roa	[x] found	docs/cognitia/agent-economy/ (7 design docs)
Latest audit repo	[x] found	docs/cognitia/audits/*, PR #188/#181/#164

Capability gaps (verified_fact from V1_1_FINAL_AUDIT.md 4 + roadmap-audit 2 + my build):

- Real channel send (email/SMS/voice/LinkedIn/ads) -- not built, by design (fence; 13).
- /gtm-command-center not on base (branch-only, 2). [x]
- CRM-2 (approval-gated stage-update) -- not built; Mira proposes only crm.task.create/crm.note.create (packages/agents/src/mira/mira.ts:174-206). [x]
- TIER-1 (risk-tiered review / earned autonomy) -- no riskTier surface in code (grep clean). [x]
- SCOPE-2 field-level write-scope check -- partial (apps/api/src/{fence,governance}.test.ts). [x]
- Production deployment, hosted/managed-Postgres RLS verification, external security audit, SOC 2 program, live pilot tenant -- all absent / founder- or infra-gated (READINESS_SCORECARD.md). [x]

6. Scorecard (honest, evidence-cited -- no axis inflated)

Sub-scores are 0-100, graded against an explicit bar with evidence. Where an axis is capped by design, that is stated.

Axis	Sco Basis / evidence	Conf
Build & CI health	92 pnpm check green, next build green on da48e8f (0)	[x]
Test breadth	82 805 tests / 106 files, two DB backends; edge-case breadth still growing	[x]
Governed CRM loop (Top-1	75 7/10 shipped + tests, 1 partial, 2 open by design (Top-10 table below)	[x]/
Alta implementation pari	70 real-module wiring on the #158 line; held below 80 -- no deploy/persist	
Dry-run automation readi	88 two tested dry-run systems + sendLive fail-closed (PR #181)	
Controlled-live readines	74 all 9 elements exist as tested code; held below the 80 gate by 4 operat	
Actual-live readiness (C	20 connector+RLS+secret store exist & conditionally wired; no deploy, no c cap	
Enterprise readiness	45 fail-closed gates/guards real; no persistence-binding, auth-binding, de	
Trust / governance desig	85 append-only proofs, doctrine guards, redaction, public diligence pack 5	[x]
Enterprise compliance ce	20 SOC 2 readiness 1/5; internal controls only; no external audit	[x]
Agent Economy (toward th	40 full loop runtime-verified but simulation-only + scope-fence divergence	
Demo / investor readines	80 2 routes build+render, deterministic run, scorecard, green tests (PR #1	[x]/
Production deployment	22 runs locally/dev only; no deployed environment (READINESS_SCORECARD.md	[x]
Pilot traction	20 product works, onboarding docs exist; no live tenant (#13 = 2/5)	[x]

Top-10 governance roadmap (roadmap-audit 2, reconciled to code/tests on da48e8f): GOV-1 [x], SIM-1 [x], TRUST-2 [x], REGR-1 [x], UNDO-1 [x], LIFE-1 [x], LEARN-1 [x] (7 shipped w/ tests); SCOPE-2 partial; CRM-2 open; TIER-1 open by design.

7. Progress bar (with the math shown)

Two honest views -- because a single number hides the deploy/sell gap.

(A) Mock-safe demonstrable MVP completeness -- mean of {Build 92, Test 82, CRM loop 75, Alta parity 70, Dry-run 88, Trust design 85, Demo 80} = $567/7 = 81\%$. ☐ 81%

(B) Production / commercial readiness -- mean of {Controlled-live 74, Actual-live 20 (capped), Enterprise 45, Compliance cert 20, Deployment 22, Pilot 20} = $201/6 = 34\%$. ☐ 34%

Headline (honest): A green, well-tested, governance-deep mock-safe MVP (~81%) that is intentionally far from deployed/sold/live (~34%). The gap is by design -- actual-live is capped -- and the dominant remaining risk is convergence, not features (8/12).

8. 7-day execution plan (convergence-first; cites the PR stack)

The roadmap-audit's lead finding: "the dominant problem is convergence, not features." Order:

1. Day 1-2 -- Land the spine toward main. Drive the merge train
#135 (w1-sales-closer-core) -> #158 (overnight/gtm-implementation) -> main. Until this lands, every Group-B PR is stranded ~3 levels deep. Adopt or explicitly supersede #149's "canonical merge/hold map". (roadmap-audit 5.1)
2. Day 2-3 -- Collapse duplicates. Pick one canonical Command Center (#168 or #169), one Proof/TrustOps PR (#165 or #166), one Alta readiness audit; close the rest. (roadmap-audit 3.5)
3. Day 3 -- Resolve the scope-fence divergence (12) with a written ruling (re-authorize + amend operating-plan.md 0a, or quarantine the /agent-economy,/credits,/skills,/proofs,/agents surfaces). This unblocks honest reporting of what is "in scope" on base.
4. Day 4-6 -- Closer automation-readiness ordered merge train (all mock-safe, all CI-green), retargeted at mainline as #158 lands: ports #171 -> dry-run ledger #175 -> controlled-live sandbox #176 -> kill-switch #174 -> consent #173 -> approval queue #177 -> monitoring #182 -> test matrix #180/#184. Release-gate #179 already in. (roadmap-audit 5.4)
5. Day 6-7 -- Wire one Command Center end-to-end so Alta parity stops being latent: drive B2 dry-run plans from the B1 assembly packet onto the B3 CRM-lite timeline, render in a real apps/web route, bind B6 permissions to the approval path (alta-80-readiness-evidence.md 3.A). mock-safe.
Gate every step on pnpm check staying green (CI: .github/workflows/ci.yml).

9. 30-day enterprise-readiness plan

Grounded in operating-plan.md 4 gates, READINESS_SCORECARD.md, and the enterprise-readiness PRs.

- Week 1: convergence (8) + scope-fence ruling.
 - Week 2 -- Persistence & deploy substrate. Provision dev/staging Postgres; apply migrations 0001-0019; verify RLS under a non-superuser role on a hosted/managed provider (the open gap in READINESS_SCORECARD.md #4/#5; plan exists at execution/MANAGED_POSTGRES_RLS_VERIFICATION_PLAN.md). Persist CRM-lite/timeline/proofs so TrustOps runs over stored data. ->[x]
 - Week 3 -- Enterprise controls bound to the route. Land #162/#185 (mock-safe enterprise-readiness: isolation, composed release decision, control matrix); bind B6 permission model to the policy/approval path; add observability/monitoring + rollback wiring (alta-80 3.B).
 - Week 4 -- Compliance program start (Gate 3 prep). SOC 2 program engaged + evidence automation (Vanta/Drata) per operating-plan.md 6; audit-trail export + retention (SEC-2); SSO/SAML spike (AUTH-2); DPA template. Compliance cert axis (6 = 20) is the lowest enterprise lever -- this is where it moves.
- Exit: a deployed, RLS-verified, monitored environment with SOC 2 Type-1 readiness engaged. No live channel send is in any exit criterion (fence, 13).

10. Agent Economy 100 plan

The Agent Economy exists today as a simulation-only loop on base (/agent-economy* API + web, work-order state machine, internal escrow, marketplace skeleton, Agent Fabric Lab) -- runtime-verified (apps/api/src/economySmoke.live.test.ts against PGLite) but with no payments, no chain, no real settlement (). Design docs: docs/cognitia/agent-economy/{WORK_ORDER_MODEL, ESCROW_SIMULATION, DISPUTE_RESOLUTION, MARKETPLACE, CROSS_TENANT_SETTLEMENT_DESIGN, AGENT_DRIVEN_WORKFLOW, AGENT_ECONOMY_LAB}.md.

Path from ~40 -> 100 (each step gated, none enabling real value transfer without explicit sign-off):

1. Resolve the scope-fence ruling first (12) -- the economy surfaces are formally outside V1 scope.
2. Cross-tenant settlement (simulated): implement CROSS_TENANT_SETTLEMENT_DESIGN.md over internal credits; keep internal_credits the only active rail (DB check-constraint enforced).
3. Dispute resolution loop: wire DISPUTE_RESOLUTION.md into the work-order state machine with proofs.
4. Marketplace matching depth: SkillProof-tier + reputation matching (PR #004/#005 branches) with listing detail pages.
5. Standards spike (no integration): map identity/reputation/settlement to ERC-8004 / EAS / x402 via the reserved external_ref columns (STANDARDS_ALIGNMENT.md) -- design only, legal gate intact.
6. Real settlement remains legal-gated (token/crypto progression in ARCHITECTURE_LOCK_V1_1.md 5); "100" is the simulated economy fully closed + standards-mapped, not a live token.

11. Investor / demo package

Ready-to-use, public-safe artifacts (no PII, no token marketing -- doctrine-guard enforced):

- Diligence overview: docs/cognitia/PUBLIC_DILIGENCE_OVERVIEW.md (2026-06-14) + Researcher Pack (docs/cognitia/public/ -- 18 docs: entypoints, FAQ, verify-it-yourself, threat model, claims-we-do-not-make, standards alignment).
 - Live demo: next build green; demo script docs/cognitia/demo/DEMO_SCRIPT_V1.md; deterministic run + scorecard; routes /cognitia, /approvals, /trust, /trust/live. Demo/investor readiness [x]/ 80-82.
 - Investor panels (branch-only): PR #161 (10 investor-grade panels) / #190 (read-only automation-readiness panel) -- pick one when the Command Center is consolidated (8).
 - Competitive narrative: docs/strategy/beat-alta-10x.md -- the accountable-revenue-action wedge into an AI-SDR category with 50-70% churn; three moats (approval-decision flywheel, in-CRM provenance, published trust benchmarks). FLY-1/PROV-1/UX-2/MET-1/EVAL-1 shipped (8 ticket table there).
 - The only non-draft PR is #89 (investor-grade audit + wedge strategy) -- surface it.
- Honesty rule for any deck: lead with the [x] axes (governed loop, tests, isolation, trust design); state the cap plainly (mock-safe, no live send, no deployed pilot). Do not claim live automation or a paying customer -- none exists in evidence (V1_1_FINAL_AUDIT.md 8; ARCHITECTURE_LOCK_V1_1.md 8).

12. Founder decision list (the human-only forks)

1. Merge-train authorization (highest leverage). Approve landing #135 -> #158 -> main and adopt/replace #149's merge/hold map. Nothing converges without this. (roadmap-audit 5.1)
2. Scope-fence ruling (4/roadmap-audit 4). operating-plan.md 0a forbids agent-economy/credits/ "Cognitia OS" code on base without written re-authorization -- yet base carries /agent-economy,/credits,/skills,/proofs,/agents (simulation-only). Re-authorize + amend 0a, or quarantine to a lab branch. As-built and as-written currently disagree.
3. Canonical Command Center pick -- #168 vs #169 (close the other ~4). (roadmap-audit 3.5)
4. Promote a default branch. Default is not the production line; READINESS_SCORECARD.md #18 flags this as a discoverability blocker. Decide overnight/gtm-implementation -> main promotion.

5. Hosted DB choice for the 9 RLS verification (e.g. Supabase/managed PG).
6. Counsel-gated items (do not start without legal): any token-shaped surface; stablecoin custody (MSB exposure); customer SMS/email consent wording (CASL) before any real send. (V1_1_FINAL_AUDIT.md 9)
7. Pilot go / no-go for a warm-network MoverOS or Demandara tenant (Tenant Map: MoverOS = Tenant Zero, Demandara = onboarding-1). (docs/cognitia/TENANT_MAP.md, pilots/DEMANDARA_PILOT_SCRIPT.md)

13. Actual-live blockers (why the cap is correct)

The platform is actual-live capped by construction -- and that cap is enforced in code, not just policy:

- packages/agents/src/channels/dryRunChannels.ts -- planDryRunAction() is a pure function that returns { mode: 'dry_run', sent: false }; assertNoLiveSend() throws on any sent:true; sendLive() always throws ("live channels disabled"). Tested in dryRunChannels.test.ts. [x]
- packages/agents/src/closer/automationReleaseGate.ts -- controlled_live **fails closed behind 7 sign-offs**; no controlled_live conditions are satisfied. [x]
- apps/api/src/killSwitch.test.ts -- any non-active integration connection halts execute/rollback (409 + audited denial). [x]
- v1 composition disables the email adapter (v1Mode=true); /webhooks/email stays disabled per operating-plan.md 0.3. [x]

Honest three-axis live posture (PR #181 live-automation-80-readiness-evidence.md,):

- Dry-run readiness 88 Controlled-live readiness 74 (below the 80 gate) Actual-live 20 (low by design; #164 scores live 12).

The 4 operational blockers holding controlled-live below 80 (PR #181): no deployed/reachable environment; post-hoc-only monitoring; the 7-signoff release gate not yet bound to the real execute path; environment unconfigured.

The blockers that must NOT be removed in any session (external sign-off only) -- alta-80 3.C / overnight-integration-plan.md 5: legal/counsel sign-off; signed customer scope + consent records; live deployment controls + connector/credential approvals; any live email/SMS/WhatsApp/call/LinkedIn/ad. Until a human clears these, actual-live stays capped -- and this report makes no live-automation or paying-customer claim.

Method & provenance

- Branch claude/cognitia-demandara-master-report-t89yho was rebuilt from origin/overnight/gtm-implementation @ da48e8f (current fetched head; confirmed the prior branch state was only the stranded vision-skill and carried no other work before reset). Diff is docs + the PDF generator only; no product code or PR state changed.
- First-party proof = the 0 commands re-run this session. All other figures are labeled (PR/branch claim), (simulated-by-design), or [?] (named input not found). Numbers cite a file/test/route/PR.
- Sources: docs/strategy/roadmap-audit-2026-06-23.md (PR #188), docs/strategy/beat-alta-10x.md, docs/competitive/operating-plan.md, docs/cognitia/{IMPLEMENTATION_COMMAND_BOOK,ARCHITECTURE_LOCK_V1_1,PUBLIC_DILIGENCE_OVERVIEW,TENANT_MAP}.md, docs/cognitia/audits/*, PRs #158/#159/#160/#164/#181/#188.